

# ATTNChecker: Highly-Optimized Fault Tolerant Attention for Large Language Model Training

Yuhang Liang<sup>1</sup>, Xinyi Li<sup>2</sup>, Jie Ren<sup>3</sup>, Ang Li<sup>2</sup>, Bo Fang<sup>2</sup>, Jieyang Chen<sup>1</sup>

<sup>1</sup>University of Oregon, <sup>2</sup>Pacific Northwest National Laboratory, <sup>3</sup>William & Mary




WILLIAM & MARY  
CHARTERED 1693

## 1. Introduction

- Training a large language model (LLM) is an extremely time and resource-intensive process, often taking weeks or even months to reach the desired accuracy.
- Recent investigations showed that the training of LLMs has high sensitivity of soft errors. The soft errors can bring the negative impact to training accuracy, even interrupt the training process.
- Current techniques, such as checkpointing, cannot efficiently handle the soft error during LLM training.
- ATTNChecker is a ABFT technique tailored for the attention mechanism in LLMs. It is designed based on fault propagation patterns of LLM and incorporates performance optimization to adapt to both system reliability and model vulnerability, while providing lightweight protection for fast LLM training.

## 2. ATTNChecker Overview

- Extreme Errors Correcting ABFT:** Handle the INF, Nan, near-INF during LLMs training, whereas the standard ABFT cannot efficiently handle these errors.
- Systematic ABFT Scheme for Attention:**
  - We divide the attention into three protection sections and allow the errors to propagate up 1D (1 column or 1 row).
  - We differentiate the **Deterministic** and **Nondeterministic** patterns and apply different strategies to handle the errors.
  - Nondeterministic:** Use both column and row checksums.
  - Deterministic:** Use single-sided checksum.
- Performance Optimizations:**
  - Customized and efficient kernels for encoding checksums.
  - Customized error detection and correction kernels for GPU to optimize the detection frequencies

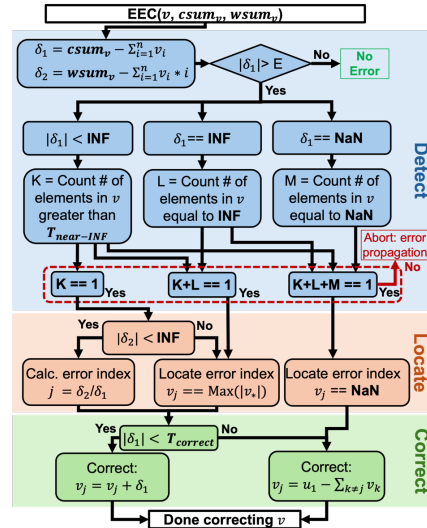


## 3. Fault Injection Study

- Experiment:** Randomly select 10% (~5,000) elements of each output matrix to inject one fault.
- Findings:**
  - Error Patterns in Attention:**
    - Errors can quickly propagate with diverse types of chances.
    - Earlier injected faults lead more negative impact to LLM training.
  - Loss:**
    - 2D pattern leads to the non-correctable state.
    - Q, K and V have higher fault sensitivity.

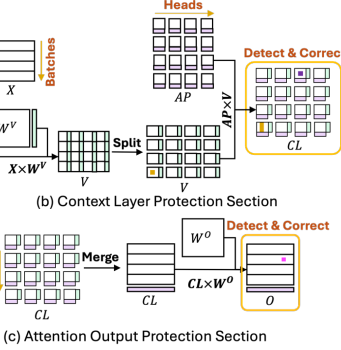

## 4. Extreme Errors Correcting ABFT

- EEC-ABFT first compares the updated checksums  $csum_v$  and  $wsum_v$  with recomputed checksum to get the difference  $\delta_1$  and  $\delta_2$  and uses  $\delta_1$  to detect the error.
- Upon detecting an error, instead of relying on a unified error handling strategy, EEC-ABFT applies corresponding error locating and correction procedures for 4 different cases by using  $\delta_1$ .



## 5. Systematic ABFT Scheme for Attention

- We divide the attention into three protection sections, shown in (a), (b) and (c).
- In each section, we allow errors propagate up to 1D and handle these 1D errors to reduce the overhead (e.g. an error occurring in Q is handled at the Attention Scores (AS)).
- Attention Scores (AS) and Context Layer (CL) are nondeterministic, as both 0D and 1D errors can occur in these matrices. To effectively handle such errors, we apply both column and row checksums.
- Attention Output (O) is deterministic, since only 0D error could occur in this matrix. Therefore, we use single-side checksum (column checksum) for error detection and correction.



## 6. Performance Optimizations

- Customized and efficient kernels for encoding checksums**

We parallelize the encoding process along the Streaming Multiprocessor by the Number of Heads  $\times$  Number of Batches to improve GPU occupancy. To encode each block, we load all data into shared memory and decoupled thread data mapping between loading and computing to achieve fully coalesced global memory accesses while minimizing bank conflict in shared memory.

  - Mapping heads to different Streaming Multiprocessors
  - Load to shared memory with coalesced patterns
  - Reordered in shared memory to minimize bank conflicts
  - Compute with hard-coded checksum weights

## Frequencies Optimization Algorithm

We treat the probability distribution of the number of error occurrences as Poisson distribution, and used a greedy algorithm, shown in Algorithm 1, to optimize the ABFT frequency.

### Algorithm 1: ABFT Frequencies Optimization

```


1 Function OptimizeABFTFrequencies( $\lambda^{INF}, \lambda^{NaN}, \lambda^{nINF}, \phi, FC_{target}$ ):
2    $f_{AS}, f_{CL}, f_{EO} \leftarrow \lambda^{INF}, \lambda^{NaN}, \lambda^{nINF}, \phi$ 
3    $FC \leftarrow 0; t_{AS}, t_{CL}, t_O \leftarrow 0$ 
4   for  $S$  in DescendSort(FCE) do
5     if  $FC_{target} - FC < FCE_S t_S$  then
6        $t_S = t_S$ 
7     end
8     if  $FC < FC_{target}$  then
9        $t_S = \frac{FC_{target} - FC}{FCE_S}$ 
10     $FC \leftarrow FC + FCE_S t_S$ 
11  end
12   $f_{AS} \leftarrow \frac{t_{AS}}{AS}; f_{CL} \leftarrow \frac{t_{CL}}{CL}; f_O \leftarrow \frac{t_O}{O}$ 
13  return  $f_{AS}, f_{CL}, f_O$ 

```


## 7. Results

### ATTNChecker Overhead

We compare ATTNChecker with the non-optimized version of ATTNChecker, which applies ABFT to all GEMMs in the attention. ATTNChecker only brings 11% overhead to attention and 7% overhead to training.




The estimated overhead of ATTNChecker during large-scale training is approximately 6.3% and remains nearly constant as the number of parameters increases.




### Recovery Overhead

Assuming the checkpointing is performed at every training step. And Checkpointing/Restore causes more than 200% overhead, while ATTNChecker achieving 49x speed up.



### Loss after 3-Epoch Training

The losses between baseline and ATTNChecker after 3-epoch training shows that ATTNChecker makes negligible impact to loss



## 8. Conclusion

We introduce ATTNChecker, an ABFT technique that offers a lightweight, real-time solution for detecting and correcting errors during LLM training. ATTNChecker is efficient in enhancing the resilience and efficiency of LLM training and provides a robust framework for future fault tolerance in large-scale machine learning models.